

# Lecture 8: GPU Programming

CSE599G1: Spring 2017


# Announcements

- **Project proposal** due on Thursday (4/28) 5pm.
- **Assignment 2** will be out today, due in two weeks.
  - Implement GPU kernels and use cublas library
  - Infer output shapes and memory planning

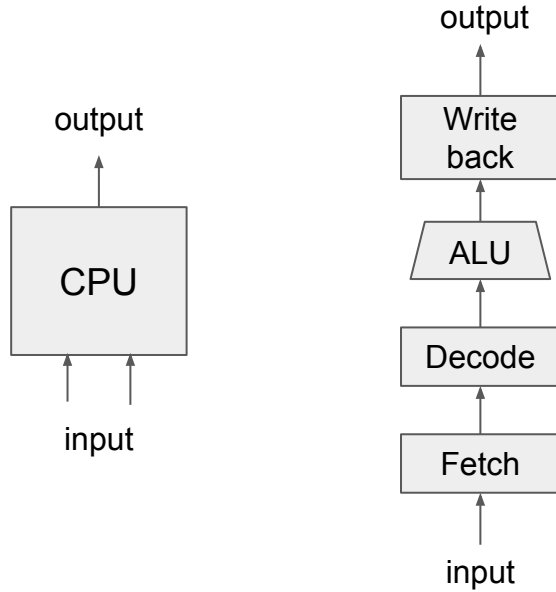
# Overview

- GPU architecture
- CUDA programming model
- Case study of efficient GPU kernels

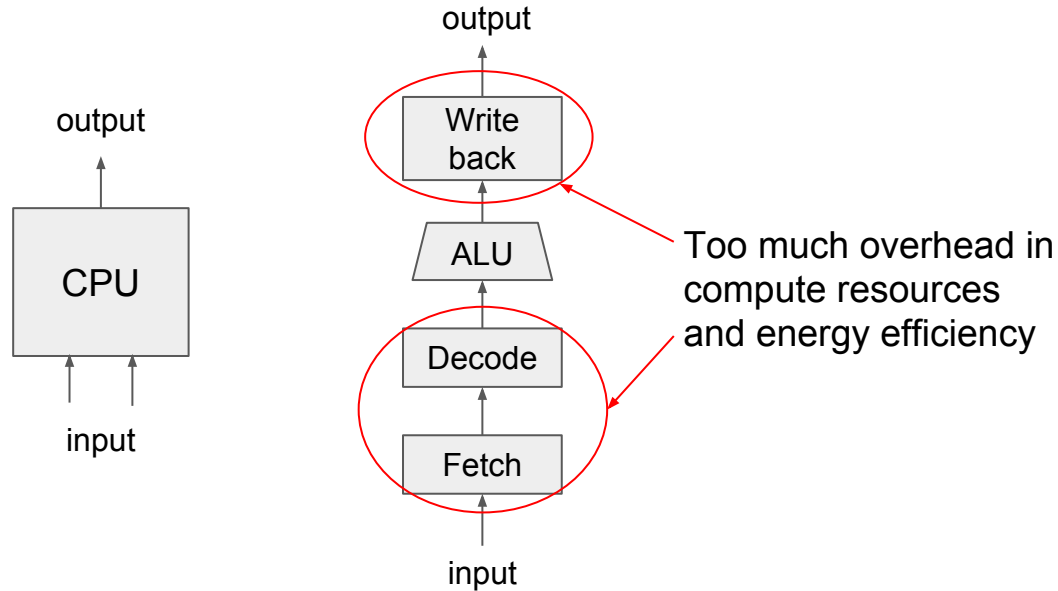
# CPU vs GPU



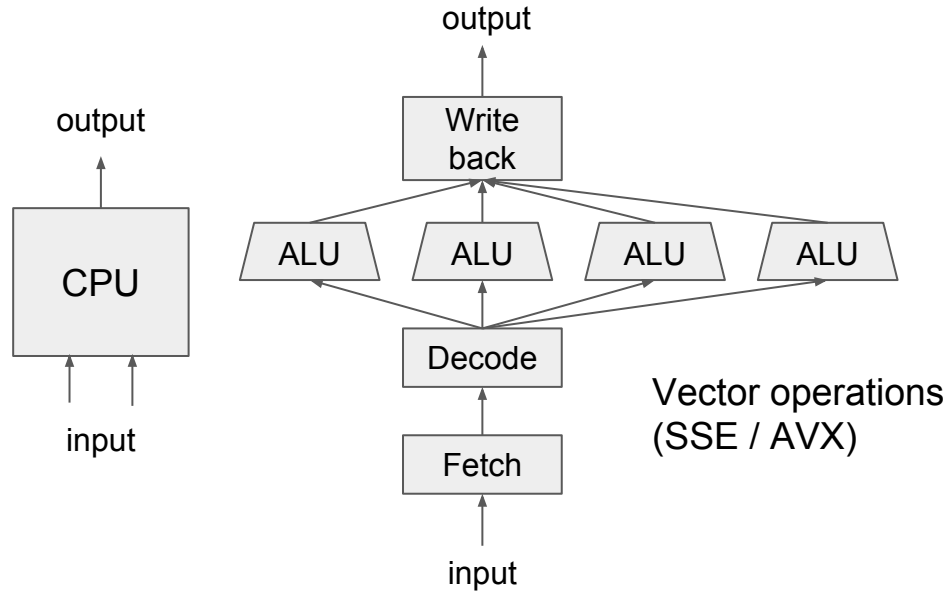
# CPU vs GPU



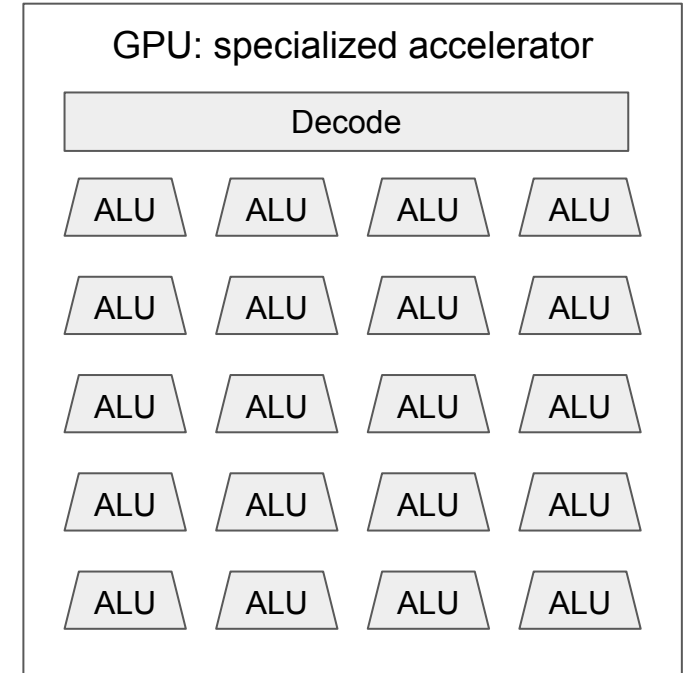
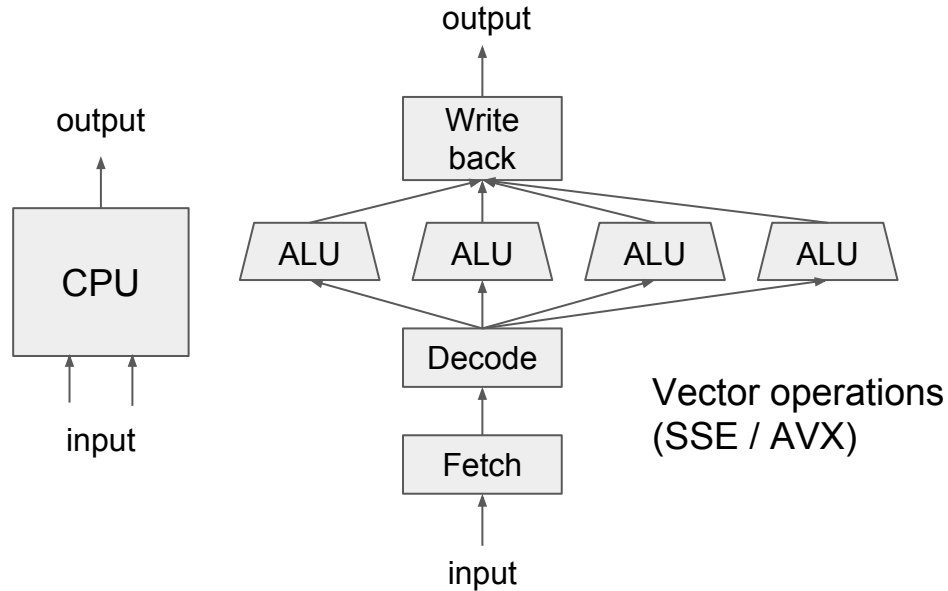
# CPU vs GPU



# CPU vs GPU



# CPU vs GPU





# Streaming Multiprocessor (SM)



Decoder: pick the next instructions

Registers

SP float core

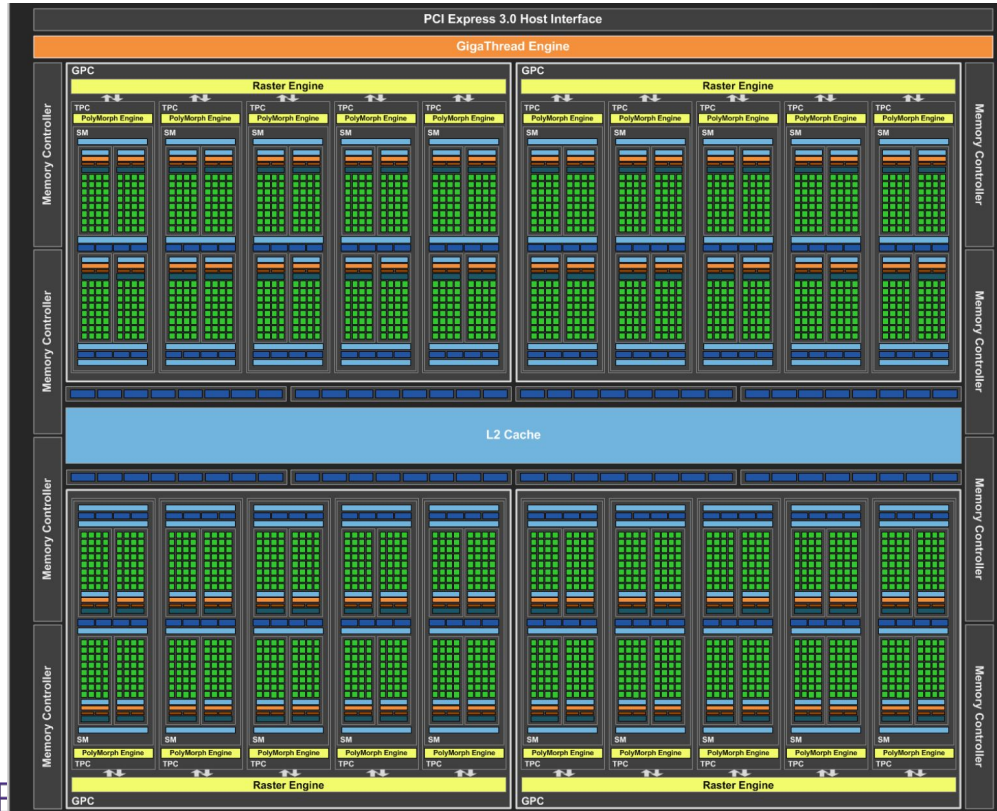
DP float core

Load/store memory


Special function unit

Multiple caches

# GPU Architecture




# Theoretical peak FLOPS comparison




# Memory Hierarchy

CPU memory hierarchy




# Memory Hierarchy

CPU memory hierarchy




# Memory Hierarchy

CPU memory hierarchy




GPU memory hierarchy

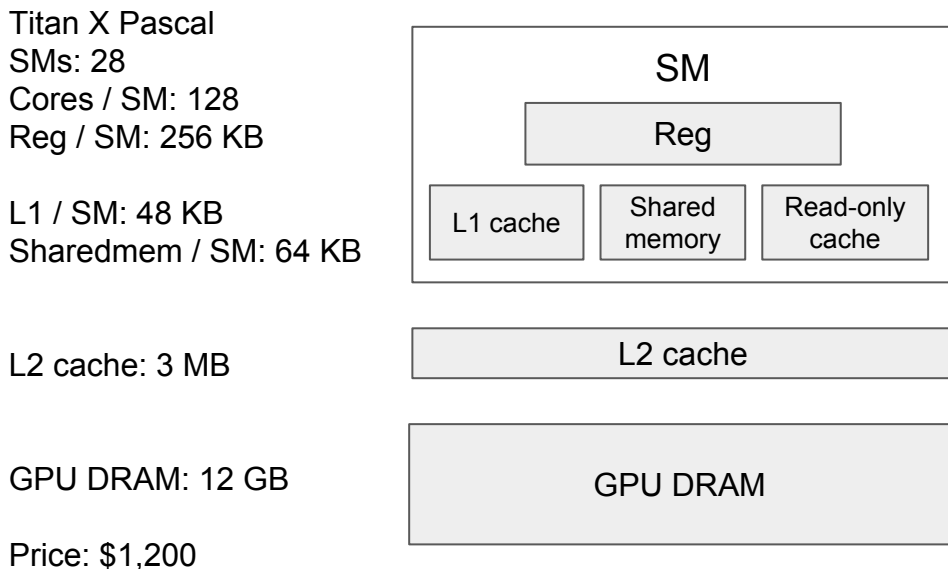


# Memory Hierarchy

## CPU memory hierarchy

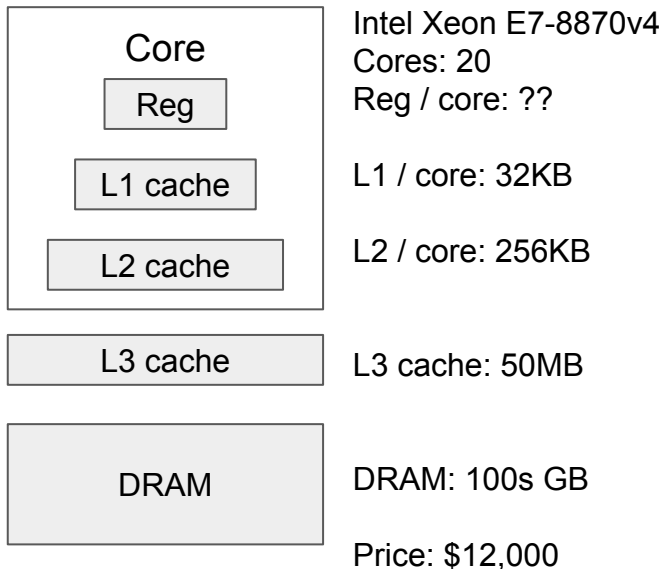


## GPU memory hierarchy




# Memory Hierarchy

## CPU memory hierarchy



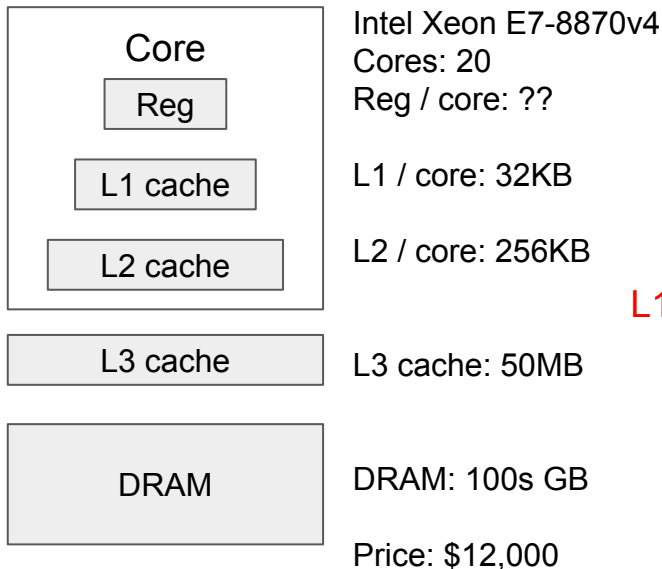
## GPU memory hierarchy



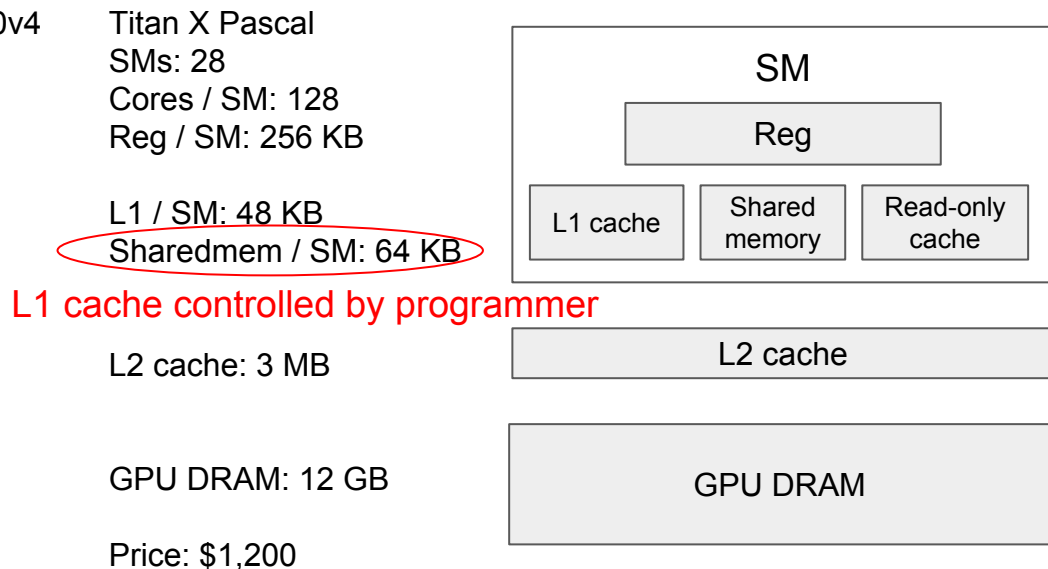


# Memory Hierarchy

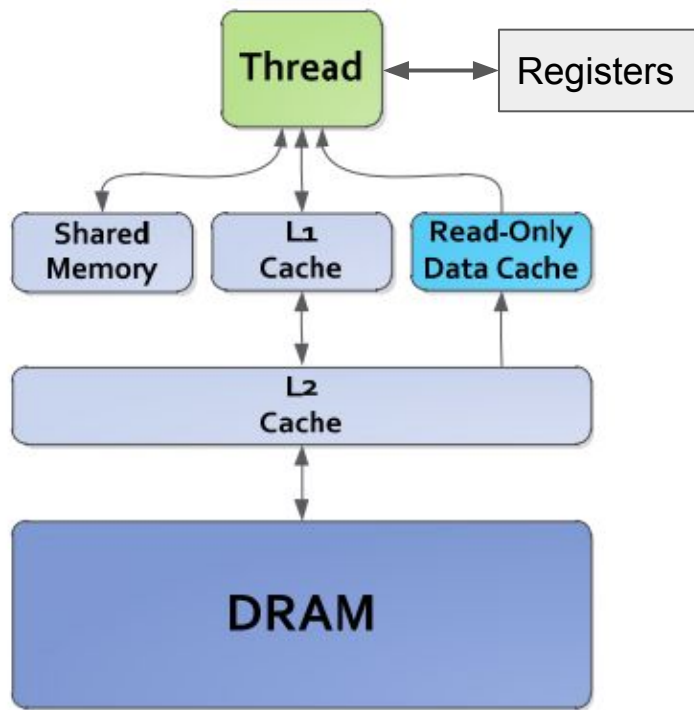
## CPU memory hierarchy



## GPU memory hierarchy



# GPU Memory Latency



Registers: R 0 cycle / R-after-W ~20 cycles

L1/texture cache: 92 cycles

Shared memory: 28 cycles

Constant L1 cache: 28 cycles


L2 cache: 200 cycles

DRAM: 350 cycles

(for Nvidia Maxwell architecture)

\* [http://lpgpu.org/wp/wp-content/uploads/2013/05/poster\\_andresch\\_acaces2014.pdf](http://lpgpu.org/wp/wp-content/uploads/2013/05/poster_andresch_acaces2014.pdf)

# Memory bandwidth comparison




# Nvidia GPU Comparison

GPU	Tesla K40 (2014)	Titan X (2015)	Titan X (2016)
Architecture	Kepler GK110	Maxwell GM200	Pascal GP102
Number of SMs	15	24	28
CUDA cores	2880 (192 * 15SM)	3072 (128 * 24SM)	3584 (128 * 28SM)
Max clock rate	875 MHz	1177 MHz	1531 MHz
FP32 GFLOPS	5040	7230	10970
32-bit Registers / SM	64K (256KB)	64K (256KB)	64K (256KB)
Shared Memory / SM	16 KB / 48 KB	96 KB	64 KB
L2 Cache / SM	1.5 MB	3 MB	3 MB
Global DRAM	12 GB	12 GB	12 GB
Power	235 W	250 W	250 W

# CUDA Programming Model

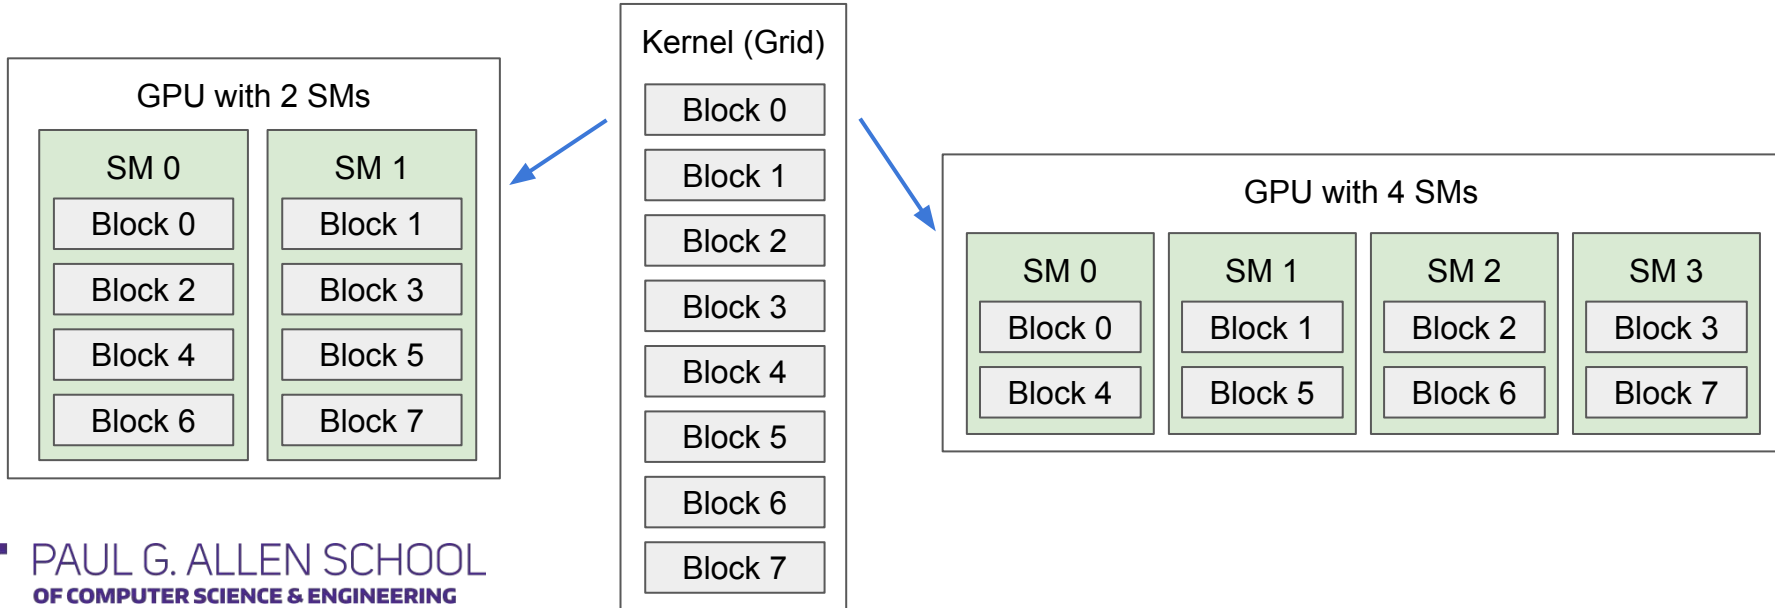
# Thread Hierarchy



- Programmer writes code for a single thread in simple C program.
  - All threads executes the same code, but can take different paths.
- Threads are grouped into a block.
  - Threads within the same block can synchronize execution.
- Blocks are grouped into a grid.
  - Blocks are independently scheduled on the GPU, can be executed in any order.
- A kernel is executed as a grid of blocks of threads.

# Kernel Execution


- Each block is executed by one SM and does not migrate.
- Several concurrent blocks can reside on one SM depending on block's memory requirement and the SM's memory resources.



# Kernel Execution




- A warp consists of 32 threads
  - A warp is the basic schedule unit in kernel execution.
- A thread block consists of 32-thread warps.
- Each cycle, a warp scheduler selects one ready warps and dispatches the warps to CUDA cores to execute.






# Thread Hierarchy & Memory Hierarchy



GPU memory hierarchy



# Example: Vector Add

```
// compute vector sum C = A + B
Void vecAdd_cpu(const float* A, const float* B, float* C, int n) {
    for (int i = 0; i < n; ++i)
        C[i] = A[i] + B[i];
}
```

# Example: Vector Add

```
// compute vector sum C = A + B
Void vecAdd_cpu(const float* A, const float* B, float* C, int n) {
    for (int i = 0; i < n; ++i)
        C[i] = A[i] + B[i];
}
```



```
__global__ void vecAddKernel(const float* A, const float* B, float* C, int n) {
    int i = blockDim.x * blockIdx.x + threadIdx.x;
    if (i < n) {
        C[i] = A[i] + B[i];
    }
}
```

# Example: Vector Add

global index	0	1	2	3	4	5	6	7	8	9	10	11
threadIdx.x	0	1	2	3	0	1	2	3	0	1	2	3
blockIdx.x	0				1				2			

Suppose each block only includes 4 threads:  
blockDim.x = 4

```
__global__ void vecAddKernel(const float* A, const float* B, float* C, int n) {  
    int i = blockDim.x * blockIdx.x + threadIdx.x;  
    if (i < n) {  
        C[i] = A[i] + B[i];  
    }  
}
```

Compute the global index

# Example: Vector Add

global index	0	1	2	3	4	5	6	7	8	9	10	11
threadIdx.x	0	1	2	3	0	1	2	3	0	1	2	3
blockIdx.x	0				1				2			

Suppose each block only includes 4 threads:  
blockDim.x = 4

```
__global__ void vecAddKernel(const float* A, const float* B, float* C, int n) {  
    int i = blockDim.x * blockIdx.x + threadIdx.x;  
    if (i < n) {  
        C[i] = A[i] + B[i];  
    }  
}
```

Each thread only performs  
one pair-wise addition

# Example: Vector Add (Host)

```
#define THREADS_PER_BLOCK    512
void vecAdd(const float* A, const float* B, float* C, int n) {
    float *d_A, *d_B, *d_C;
    int size = n * sizeof(float);
    cudaMalloc((void **) &d_A, size);
    cudaMemcpy(d_A, A, size, cudaMemcpyHostToDevice);
    cudaMalloc((void **) &d_B, size);
    cudaMemcpy(d_B, B, size, cudaMemcpyHostToDevice);
    cudaMalloc((void **) &d_C, size);
    int nblocks = (n + THREADS_PER_BLOCK - 1) / THREADS_PER_BLOCK;
    vecAddKernel<<<nblocks, THREADS_PER_BLOCK>>>(d_A, d_B, d_C, n);
    cudaMemcpy(C, d_C, size, cudaMemcpyDeviceToHost);
    cudaFree(d_A); cudaFree(d_B); cudaFree(d_C);
}
```

# Example: Vector Add (Host)


```
#define THREADS_PER_BLOCK    512
void vecAdd(const float* A, const float* B, float* C, int n) {
    float *d_A, *d_B, *d_C;
    int size = n * sizeof(float);
    cudaMalloc((void **) &d_A, size);
    cudaMemcpy(d_A, A, size, cudaMemcpyHostToDevice);
    cudaMalloc((void **) &d_B, size);
    cudaMemcpy(d_B, B, size, cudaMemcpyHostToDevice);
    cudaMalloc((void **) &d_C, size);
    int nblocks = (n + THREADS_PER_BLOCK - 1) / THREADS_PER_BLOCK;
    vecAddKernel<<<nblocks, THREADS_PER_BLOCK>>>(d_A, d_B, d_C, n);
    cudaMemcpy(C, d_C, size, cudaMemcpyDeviceToHost);
    cudaFree(d_A); cudaFree(d_B); cudaFree(d_C);
}
```

Launch the GPU kernel  
asynchronously

# Example: Sliding Window Sum

- Consider computing the sum of a sliding window over a vector
  - Each output element is the sum of input elements within a radius
  - Example: image blur kernel
- If radius is 3, each output element is sum of 7 input elements

input



output





# A naive implementation

```
#define RADIUS 3
__global__ void windowSumNaiveKernel(const float* A, float* B, int n) {
    int out_index = blockDim.x * blockIdx.x + threadIdx.x;
    int in_index = out_index + RADIUS;
    if (out_index < n) {
        float sum = 0.;
        for (int i = -RADIUS; i <= RADIUS; ++i) {
            sum += A[in_index + i];
        }
        B[out_index] = sum;
    }
}
```

# A naive implementation


```
void windowSum(const float* A, float* B, int n) {  
    float *d_A, *d_B;  
    int size = n * sizeof(float);  
    cudaMalloc((void **) &d_A, (n + 2 * RADIUS) * sizeof(float));  
    cudaMemset(d_A, 0, (n + 2 * RADIUS) * sizeof(float));  
    cudaMemcpy(d_A + RADIUS, A, size, cudaMemcpyHostToDevice);  
    cudaMalloc((void **) &d_B, size);  
    dim3 threads(THREADS_PER_BLOCK, 1, 1);  
    dim3 blocks((n + THREADS_PER_BLOCK - 1) / THREADS_PER_BLOCK, 1, 1);  
    windowSumNaiveKernel<<<blocks, threads>>>(d_A, d_B, n);  
    cudaMemcpy(B, d_B, size, cudaMemcpyDeviceToHost);  
    cudaFree(d_A); cudaFree(d_B);  
}
```

# How to improve it?

- For each element in the input, how many times it is loaded?

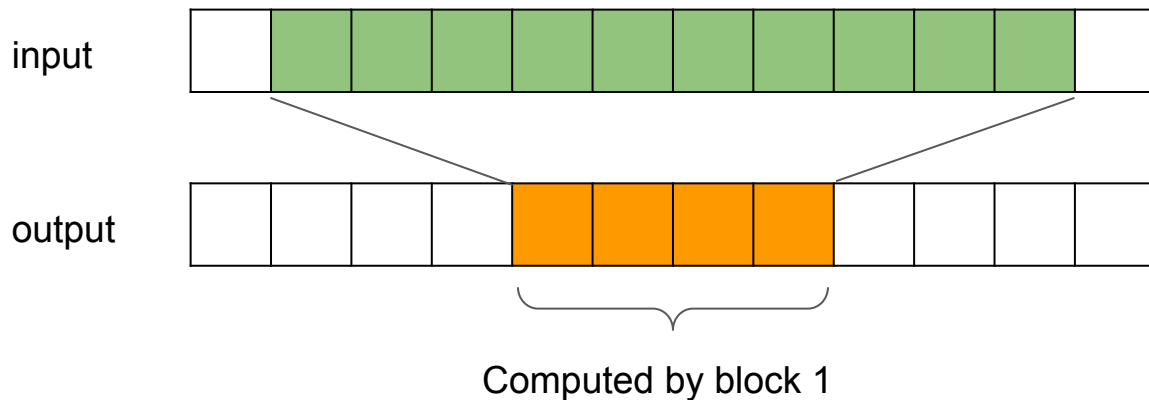
# How to improve it?

- For each element in the input, how many times it is read?
  - Each input element is read 7 times!
  - Neighboring threads read most of the same elements
- How can we avoid redundant reading of data?



# Sharing data between threads within a block

- A thread block first cooperatively loads the needed input data into the shared memory.



# Kernel with shared memory

```
__global__ void windowSumKernel(const float* A, float* B, int n) {  
    __shared__ float temp[THREADS_PER_BLOCK + 2 * RADIUS];  
    int out_index = blockDim.x * blockIdx.x + threadIdx.x;  
    int in_index = out_index + RADIUS;  
    int local_index = threadIdx.x + RADIUS;  
    if (out_index < n) {  
        temp[local_index] = A[in_index];  
        if (threadIdx.x < RADIUS) {  
            temp[local_index - RADIUS] = A[in_index - RADIUS];  
            temp[local_index + THREADS_PER_BLOCK] = A[in_index+THREADS_PER_BLOCK];  
        }  
        __syncthreads();  
    }  
}
```

# Kernel with shared memory

```
float sum = 0.;  
for (int i = -RADIUS; i <= RADIUS; ++i) {  
    sum += temp[local_index + i];  
}  
B[out_index] = sum;  
}  
}
```

# Performance comparison

Demo!

Code:


[https://github.com/dlsys-course/dlsys-course.github.io/blob/master/examples/window\\_sum.cu](https://github.com/dlsys-course/dlsys-course.github.io/blob/master/examples/window_sum.cu)



# Case study of efficient GPU kernels

# Case study: GEMM


$C = A \times B$   
A:  $M \times K$  matrix  
B:  $K \times N$  matrix  
C:  $M \times N$  matrix



Workload of a thread block

# Case study: GEMM

$C = A \times B$   
A:  $M \times K$  matrix  
B:  $K \times N$  matrix  
C:  $M \times N$  matrix



Workload of a thread block


# Case study: GEMM

$$C = A \times B$$

A:  $M \times K$  matrix

B:  $K \times N$  matrix

C:  $M \times N$  matrix



Workload of a thread block


# Case study: GEMM

$$C = A \times B$$

A:  $M \times K$  matrix

B:  $K \times N$  matrix

C:  $M \times N$  matrix



# Case study: GEMM pseudocode

block\_dim:  $\langle M / b, N / b \rangle$

thread\_dim:  $\langle t, t \rangle$

// thread function

```
__global__ void SGEMM(float *A, float *B, float *C, int b, int s) {  
    __shared__ float sA[2][b][s], sB[2][s][b]; // shared by a thread block  
    float rC[bt][bt] = {0}; // thread local buffer, in the registers  
    Cooperative fetch first strip from A, B to sA[0], sB[0]  
    __sync_threads();  
    for (k = 0; k < K / s; k += 1) {  
        Cooperative fetch next strip from A, B to sA[(k+1)%2], sB[(k+1)%2]  
        __sync_threads();  
        for (kk = 0; kk < s; kk += 1) {  
            for (j = 0; j < bt; j += 1) { // unroll loop  
                for (i = 0; i < bt; i += 1) { // unroll loop  
                    rC[j][i] += sA[k%2][threadIdx.x*bt+j][kk]*sB[k%2][kk][threadIdx.y*bt+i];  
                }  
            }  
        }  
    }  
}
```

Run in parallel



Write rC back to C

# Case study: GEMM

More details see:

- <http://homes.cs.washington.edu/~tw510/cse599i/CSE%20599%20I%20Accelerated%20Computing%20-%20Programming%20GPUs%20Lecture%204.pdf>
- Lai, Junjie, and André Seznec. "Performance upper bound analysis and optimization of SGEMM on Fermi and Kepler GPUs." Code Generation and Optimization (CGO), 2013 IEEE/ACM International Symposium on. IEEE, 2013.

# Case study: Reduction Sum

[http://developer.download.nvidia.com/compute/cuda/1.1-Beta/x86\\_website/projects/reduction/doc/reduction.pdf](http://developer.download.nvidia.com/compute/cuda/1.1-Beta/x86_website/projects/reduction/doc/reduction.pdf)



# Tips for high performance

- Use existing libraries, which are highly optimized, e.g. cublas, cudnn.
- Use nvprof or nvvp (visual profiler) to debug the performance.
- Use high level language to write GPU kernels.

# References

- CUDA Programming Guide:  
<http://docs.nvidia.com/cuda/cuda-c-programming-guide/>